

# Deep Learning

## Chapter 1: A Single Neuron

David Rügamer

Department of Statistics – LMU Munich

Winter Semester 2020



# A SINGLE NEURON

- In order to better understand the types of functions that neural networks can represent, let us begin with a very simple model: logistic regression.
- Recall: The hypothesis space of logistic regression can be written as

$$\mathcal{H} = \left\{ f : \mathbb{R}^p \rightarrow [0, 1] \mid f(\mathbf{x}) = \tau \left( \sum_{j=1}^p w_j x_j + b \right), \mathbf{w} \in \mathbb{R}^p, b \in \mathbb{R} \right\},$$

where  $\tau(z) = (1 + \exp(-z))^{-1}$  is the logistic sigmoid function.

- It is straightforward to represent this function  $f(\mathbf{x})$  graphically as a neuron.
- Note:  $\mathbf{w}$  and  $b$  together constitute  $\theta$ .

# A SINGLE NEURON

We consider a logistic regression model for  $p = 3$ , i.e.

$$f(\mathbf{x}) = \tau(w_1 \mathbf{x}_1 + w_2 \mathbf{x}_2 + w_3 \mathbf{x}_3 + b).$$

- First, features of  $\mathbf{x}$  are represented by nodes in the “input layer”.

Input




- In general, a  $p$ -dimensional input vector  $\mathbf{x}$  will be represented by  $p$  nodes in the input layer.

# A SINGLE NEURON

We consider a logistic regression model for  $p = 3$ , i.e.

$$f(\mathbf{x}) = \tau(w_1 x_1 + w_2 x_2 + w_3 x_3 + b).$$

- Next, weights  $\mathbf{w}$  are represented by edges from the input layer.




- The bias term  $b$  is implicit. It is not shown/represented visually.

# A SINGLE NEURON

**Note:** For an explicit graphical representation of the bias term, we can do a simple trick:

- we add a constant feature to the inputs  $\tilde{\mathbf{x}} = (1, x_1, \dots, x_p)^\top$
- and the bias term to the weight vector  $\tilde{\mathbf{w}} = (b, w_1, \dots, w_p)$ .

The graphical representation is then:




**Figure:** Weights and bias of the neuron.

# A SINGLE NEURON

We consider a logistic regression model for  $p = 3$ , i.e.

$$f(\mathbf{x}) = \tau(w_1 x_1 + w_2 x_2 + w_3 x_3 + b).$$

- Finally, the computation  $\tau(w_1 x_1 + w_2 x_2 + w_3 x_3 + b)$  is represented by the neuron in the “output layer”.




Because this single neuron represents exactly the same hypothesis space as logistic regression, it can only learn linear decision boundaries.

# A SINGLE NEURON

- Therefore, a neuron is just a graphical representation of a very specific kind of function.
- Every neuron performs a 2-step computation:
  - Step 1: compute the weighted sum of inputs (with bias).
  - Step 2: apply an **activation function** to the sum, which is usually a non-linear transformation of the input.
- With a single neuron, the activation function serves to constrain the output to the desired range of values. In case of logistic regression. For example, it squashes the output into  $[0, 1]$ .
- However, we will see that activation functions serve a far more important purpose. They are one of the main reasons that neural networks can represent extremely complicated functions.

# A SINGLE NEURON

- One very nice thing about the graphical representation of the function is that you can picture the input vector being "fed" to the neuron on the left followed by a sequence of computations being performed from left to right. This is called a "**forward pass**".




**Figure:** Weights (and bias) of the neuron.



# A SINGLE NEURON


- One very nice thing about the graphical representation of the function is that you can picture the input vector being "fed" to the neuron on the left followed by a sequence of computations being performed from left to right. This is called a "**forward pass**".



**Figure:** Feed the input on the left.

# A SINGLE NEURON

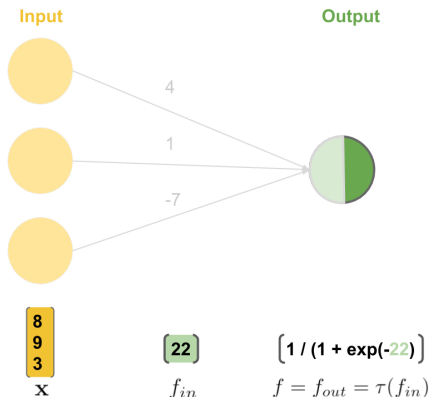
- One very nice thing about the graphical representation of the function is that you can picture the input vector being "fed" to the neuron on the left followed by a sequence of computations being performed from left to right. This is called a "**forward pass**".



**Figure:** Step 1: Compute the weighted sum.

# A SINGLE NEURON


- One very nice thing about the graphical representation of the function is that you can picture the input vector being "fed" to the neuron on the left followed by a sequence of computations being performed from left to right. This is called a "**forward pass**".



**Figure:** Step 2: Apply the activation function.


# A SINGLE NEURON

- One very nice thing about the graphical representation of the function is that you can picture the input vector being "fed" to the neuron on the left followed by a sequence of computations being performed from left to right. This is called a "**forward pass**".



# A SINGLE NEURON

- Even though all neurons compute a weighted sum in the first step, there is considerable flexibility in the type of activation function used in the second step.
- For example, setting the activation function to the identity function allows a neuron to represent linear regression.




# A SINGLE NEURON

- The hypothesis space that is formed by single neuron architectures is

$$\mathcal{H} = \left\{ f : \mathbb{R}^p \rightarrow \mathbb{R} \mid f(\mathbf{x}) = \tau \left( \sum_{j=1}^p w_j x_j + b \right), \mathbf{w} \in \mathbb{R}^p, b \in \mathbb{R} \right\}.$$

- Both logistic regression and linear regression are subspaces of  $\mathcal{H}$  (if  $\tau$  is the logistic sigmoid / identity function).



**Figure:** *Left:* A regression line learned by a single neuron. *Right:* A decision-boundary learned by a single neuron in a binary classification task.

# A SINGLE NEURON: OPTIMIZATION

- To optimize this model, we minimize the empirical risk

$$\mathcal{R}_{\text{emp}} = \sum_{i=1}^n L\left(y^{(i)}, f\left(\mathbf{x}^{(i)}\right)\right),$$

where  $L(y, f(\mathbf{x}))$  is a loss function. It compares the network's predictions  $f(\mathbf{x})$  to the ground truth  $y$ .

- For regression, we typically use the L2 loss (rarely L1):

$$L(y, f(\mathbf{x})) = \frac{1}{2}(y - f(\mathbf{x}))^2$$

- For binary classification, we typically apply the cross entropy loss (also known as bernoulli loss):

$$L(y, f(\mathbf{x})) = y \log f(\mathbf{x}) + (1 - y) \log(1 - f(\mathbf{x}))$$

# A SINGLE NEURON: OPTIMIZATION

- For a single neuron, in both cases, the loss function is convex and the global optimum can be found with an iterative algorithm like gradient descent.
- In fact, a single neuron with logistic sigmoid function trained with the bernoulli loss does not only have the same hypothesis space as a logistic regression and is therefore the same model, but will also yield to the very same result when trained until convergence.