

Deep Learning

Convolutional Operation

a	b	c
d	e	f
g	h	i

Input: 3x3x1

w_{11}	w_{12}
w_{21}	w_{22}

Filter: 2x2x1

s_{11}	s_{12}
s_{21}	s_{22}

Output: 2x2x1

$$s_{11} = a \cdot w_{11} + b \cdot w_{12} + d \cdot w_{21} + e \cdot w_{22}$$

$$s_{12} = b \cdot w_{11} + c \cdot w_{12} + e \cdot w_{21} + f \cdot w_{22}$$

$$s_{21} = d \cdot w_{11} + e \cdot w_{12} + g \cdot w_{21} + h \cdot w_{22}$$

$$s_{22} = e \cdot w_{11} + f \cdot w_{12} + h \cdot w_{21} + i \cdot w_{22}$$

Learning goals

- What are the filters?
- Convolutional Operation
- 2D Convolution

FILTERS TO EXTRACT FEATURES

- Filters are widely applied in Computer Vision (CV) since the 70's.
- One prominent example: **Sobel-Filter**.
- It detects edges in images.




Figure: Sobel-filtered image.

FILTERS TO EXTRACT FEATURES

- Edges occur where the intensity over neighboring pixels changes fast.
- Thus, approximate the gradient of the intensity of each pixel.
- Sobel showed that the gradient image \mathbf{G}_x of original image \mathbf{A} in x-dimension can be approximated by:

$$\mathbf{G}_x = \begin{bmatrix} -1 & 0 & +1 \\ -2 & 0 & +2 \\ -1 & 0 & +1 \end{bmatrix} * \mathbf{A} = \mathbf{S}_x * \mathbf{A}$$

where $*$ indicates a mathematical operation known as a **convolution**, not a traditional matrix multiplication.

- The filter matrix \mathbf{S}_x consists of the product of an **averaging** and a **differentiation** kernel:

$$\underbrace{\begin{bmatrix} 1 & 2 & 1 \end{bmatrix}}_{\text{averaging}}^T \underbrace{\begin{bmatrix} -1 & 0 & +1 \end{bmatrix}}_{\text{differentiation}}$$

FILTERS TO EXTRACT FEATURES

- Similarly, the gradient image \mathbf{G}_y in y-dimension can be approximated by:

$$\mathbf{G}_y = \begin{bmatrix} -1 & -2 & -1 \\ 0 & 0 & 0 \\ +1 & +2 & +1 \end{bmatrix} * \mathbf{A} = \mathbf{S}_y * \mathbf{A}$$

- The combination of both gradient images yields a dimension-independent gradient information \mathbf{G} :

$$\mathbf{G} = \sqrt{\mathbf{G}_x^2 + \mathbf{G}_y^2}$$

- These matrix operations were used to create the filtered picture of Albert Einstein.


HORIZONTAL VS VERTICAL EDGES



Input



Vertical edges detected by S_x



Horizontal edges detected by S_y

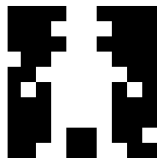


Combined

Source: Wikipedia

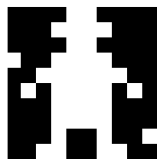
Figure: Sobel filtered images. Outputs are normalized in each case.

FILTERS TO EXTRACT FEATURES



- Let's do this on a dummy image.
- How to represent a digital image?

FILTERS TO EXTRACT FEATURES



0	0	0	0	255	255	0	0	0	0
0	0	0	255	255	255	255	0	0	0
0	0	0	0	255	255	0	0	0	0
255	0	0	255	255	255	255	0	0	0
0	0	255	255	255	255	255	255	0	0
0	255	0	255	255	255	255	0	255	0
0	0	0	255	255	255	255	0	0	0
0	0	0	255	255	255	255	0	0	0
0	0	0	255	255	255	255	0	0	0
0	0	0	255	0	0	255	0	0	255
0	0	255	255	0	0	255	255	0	0

- Basically as an array of integers.

FILTERS TO EXTRACT FEATURES

Sobel-Operator

$$S_x = \begin{bmatrix} -1 & 0 & +1 \\ -2 & \boxed{0} & +2 \\ -1 & 0 & +1 \end{bmatrix}$$

0	0	0	0	255	255	0	0	0	0
0	0	0	255	255	255	255	0	0	0
0	0	0	0	255	255	0	0	0	0
255	0	0	255	255	255	255	0	0	0
0	0	255	255	255	255	255	255	0	0
0	255	0	255	255	255	255	0	255	0
0	0	0	255	255	255	255	0	0	0
0	0	0	255	255	255	255	0	0	0
0	0	0	255	0	0	255	0	0	255
0	0	255	255	0	0	255	255	0	0

- S_x enables us to detect vertical edges!

FILTERS TO EXTRACT FEATURES

Sobel-Operator

$$S_x = \begin{bmatrix} -1 & 0 & +1 \\ -2 & 0 & +2 \\ -1 & 0 & +1 \end{bmatrix}$$

0	0	0	0	255	255	0	0	0	0
0	0	0	255	255	255	255	0	0	0
0	0	0	0	0	255	255	0	0	0
255	0	0	255	255	255	255	255	0	0
0	0	255	255	255	255	255	255	0	0
0	255	0	255	255	255	255	0	255	0
0	0	0	255	255	255	255	0	0	0
0	0	0	255	255	255	255	0	0	0
0	0	0	255	0	0	255	0	0	255
0	0	255	255	0	0	255	255	0	0

FILTERS TO EXTRACT FEATURES

Sobel-Operator

$$S_X = \begin{bmatrix} -1 & 0 & +1 \\ -2 & \boxed{0} & +2 \\ -1 & 0 & +1 \end{bmatrix}$$

0	0	0	0	255	255	0	0	0	0
0	0	0	255	255	255	255	0	0	0
0	0	0	0	255	255	255	0	0	0
255	0	0	255	255	255	255	255	0	0
0	0	255	255	255	255	255	255	0	0
0	255	0	255	255	255	255	0	255	0
0	0	0	255	255	255	255	0	0	0
0	0	0	255	255	255	255	0	0	0
0	0	0	255	0	0	255	0	0	255
0	0	255	255	0	0	255	255	0	0

$$\begin{aligned} (G_X)_{(i,j)} &= (I \star S_X)_{(i,j)} = -1 \cdot 0 + 0 \cdot 255 + \mathbf{1 \cdot 255} \\ &\quad - 2 \cdot 0 + 0 \cdot 0 + \mathbf{2 \cdot 255} \\ &\quad - 1 \cdot 0 + 0 \cdot 255 + \mathbf{1 \cdot 255} \end{aligned}$$

FILTERS TO EXTRACT FEATURES

0	510	1020	510	-510	-1020	-510	0
-255	510	1020	510	-510	-1020	-510	0
-255	765	765	255	-255	-765	-765	-255
255	765	510	0	0	-510	-765	-510
255	510	765	0	0	-765	-510	-255
0	765	1020	0	0	-1020	-765	0
0	1020	765	-255	255	-765	-1020	255
255	1020	0	-765	765	0	-1020	255

- Applying the Sobel-Operator to every location in the input yields us the **feature map**.

FILTERS TO EXTRACT FEATURES



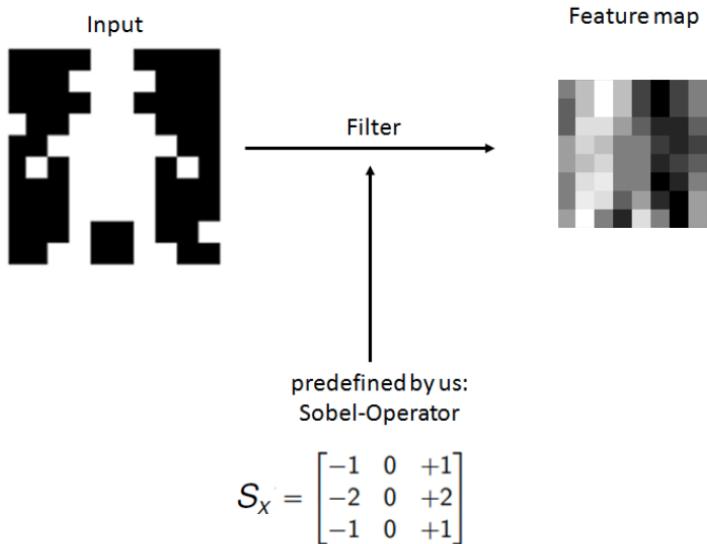
128	191	255	191	64	0	64	128
96	191	255	191	64	0	64	128
96	223	223	159	96	32	32	96
159	223	191	128	128	64	32	64
159	191	223	128	128	32	64	96
128	223	255	128	128	0	32	128
128	255	223	96	159	32	0	159
159	255	128	32	223	128	0	159

- Normalized feature map reveals vertical edges.
- Note the dimensional reduction compared to the dummy image.

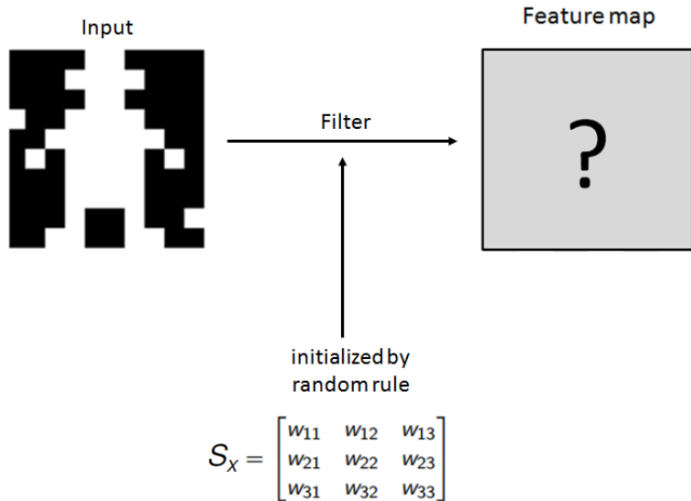
WHY DO WE NEED TO KNOW ALL OF THAT?

- What we just did was extracting **pre-defined** features from our input (i.e. edges).
- A convolutional neural network does almost exactly the same: “extracting features from the input”.
⇒ The main difference is that we usually do not tell the CNN what to look for (pre-define them), **the CNN decides itself**.
- In a nutshell:
 - We initialize a lot of random filters (like the Sobel but just random entries) and apply them to our input.
 - Then, a classifier which is generally a feed forward neural net, uses them as input data.
 - Filter entries will be adjusted by common gradient descent methods.

WHY DO WE NEED TO KNOW ALL OF THAT?



WHY DO WE NEED TO KNOW ALL OF THAT?



WORKING WITH IMAGES

- In order to understand the functionality of CNNs, we have to familiarize ourselves with some properties of images.
- Grey scale images:
 - Matrix with dimensions **height** \times **width** \times 1.
 - Pixel entries differ from 0 (black) to 255 (white).
- Color images:
 - Tensor with dimensions **height** \times **width** \times 3.
 - The depth 3 denotes the RGB values (red - green - blue).
- Filters:
 - A filter's depth is **always** equal to the input's depth!
 - In practice, filters are usually square.
 - Thus we only need one integer to define its size.
 - For example, a filter of size 2 applied on a color image actually has the dimensions $2 \times 2 \times 3$.

THE 2D CONVOLUTION

- Suppose we have an input with entries a, b, \dots, i (think of pixel values).
- The filter we would like to apply has weights w_{11}, w_{12}, w_{21} and w_{22} .

a	b	c
d	e	f
g	h	i

Input: 3x3x1

w_{11}	w_{12}
w_{21}	w_{22}

Filter: 2x2x1

THE 2D CONVOLUTION

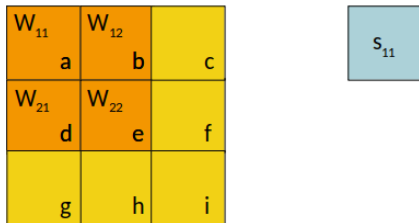
- Suppose we have an input with entries a, b, \dots, i (think of pixel values).
- The filter we would like to apply has weights w_{11}, w_{12}, w_{21} and w_{22} .

w_{11} a	w_{12} b	c
w_{21} d	w_{22} e	f
g	h	i



THE 2D CONVOLUTION

- Suppose we have an input with entries a, b, \dots, i (think of pixel values).
- The filter we would like to apply has weights w_{11}, w_{12}, w_{21} and w_{22} .



To obtain s_{11} we simply compute the dot product:

$$s_{11} = a \cdot w_{11} + b \cdot w_{12} + d \cdot w_{21} + e \cdot w_{22}$$

THE 2D CONVOLUTION

- Suppose we have an input with entries a, b, \dots, i (think of pixel values).
- The filter we would like to apply has weights w_{11}, w_{12}, w_{21} and w_{22} .

	w_{11}	w_{12}
a	b	c
d	e	f
g	h	i

s_{11}	s_{12}
----------	----------

Same for s_{12} :

$$s_{12} = b \cdot w_{11} + c \cdot w_{12} + e \cdot w_{21} + f \cdot w_{22}$$

THE 2D CONVOLUTION

- Suppose we have an input with entries a, b, \dots, i (think of pixel values).
- The filter we would like to apply has weights w_{11}, w_{12}, w_{21} and w_{22} .

	a	b	c
w_{11}	d	w_{12}	e
w_{21}	g	w_{22}	h

s_{11}	s_{12}
s_{21}	

As well as for s_{21} :

$$s_{21} = d \cdot w_{11} + e \cdot w_{12} + g \cdot w_{21} + h \cdot w_{22}$$

THE 2D CONVOLUTION

- Suppose we have an input with entries a, b, \dots, i (think of pixel values).
- The filter we would like to apply has weights w_{11}, w_{12}, w_{21} and w_{22} .

	a	b	c
d	w_{11}	w_{12}	e
g	w_{21}	w_{22}	h

s_{11}	s_{12}
s_{21}	s_{22}

And finally for s_{22} :

$$s_{22} = e \cdot w_{11} + f \cdot w_{12} + h \cdot w_{21} + i \cdot w_{22}$$

THE 2D CONVOLUTION

- Suppose we have an input with entries a, b, \dots, i (think of pixel values).
- The filter we would like to apply has weights w_{11}, w_{12}, w_{21} and w_{22} .

a	b	c
d	e	f
g	h	i

Input: 3x3x1

w_{11}	w_{12}
w_{21}	w_{22}

Filter: 2x2x1

s_{11}	s_{12}
s_{21}	s_{22}

Output: 2x2x1

$$s_{11} = a \cdot w_{11} + b \cdot w_{12} + d \cdot w_{21} + e \cdot w_{22}$$

$$s_{12} = b \cdot w_{11} + c \cdot w_{12} + e \cdot w_{21} + f \cdot w_{22}$$

$$s_{21} = d \cdot w_{11} + e \cdot w_{12} + g \cdot w_{21} + h \cdot w_{22}$$

$$s_{22} = e \cdot w_{11} + f \cdot w_{12} + h \cdot w_{21} + i \cdot w_{22}$$

THE 2D CONVOLUTION

- Suppose we have an input with entries a, b, \dots, i (think of pixel values).
- The filter we would like to apply has weights w_{11}, w_{12}, w_{21} and w_{22} .

a	b	c
d	e	f
g	h	i

Input: 3x3x1

w_{11}	w_{12}
w_{21}	w_{22}

Filter: 2x2x1

s_{11}	s_{12}
s_{21}	s_{22}

Output: 2x2x1

More generally, let I be the matrix representing the input and W be the filter/kernel. Then the entries of the output matrix are defined by $s_{ij} = \sum_{m,n} I_{i+m-1,j+n-1} w_{mn}$ where m, n denote the image size and kernel size respectively.