Deep Learning

Attention and Transformers

A woman is throwing a frisbee in a park.

A stop sign is on a road with a mountain in the background.

Learning goals

- Familiarize with the most recent sequence data modeling technique:
  - Attention Mechanism
  - Transformers
- Get to know the CNN alternative to RNNs

Attention

WAHT IS ATTENTION

Humans process data by actively shifting their focus:

- Different parts of an image carry different information
- Words derive their specific meaning from **contex**
- Remember specific, **related** events in the past
- Allows to follow one thought at a time while suppressing information **irrelevant** to the task
- Example: cocktail party problem

WAHT IS ATTENTION

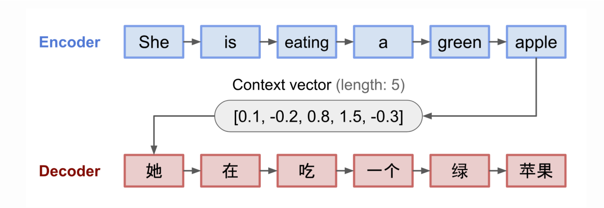

Figure: The encoder-decoder model for translation. (source:https://lilianweng.github.io/lil-log/2018/06/24/attention-attention.html)

- In a classical decoder-encoder RNN all information about the input sequence must be incorporated into the final hidden state, which is then passed as an input to the decoder network.
- With a long input sequence this fixed-sized context vector is unlikely to capture all relevant information about the past.
- Each hidden state contains mostly information from recent inputs.

WAHT IS ATTENTION

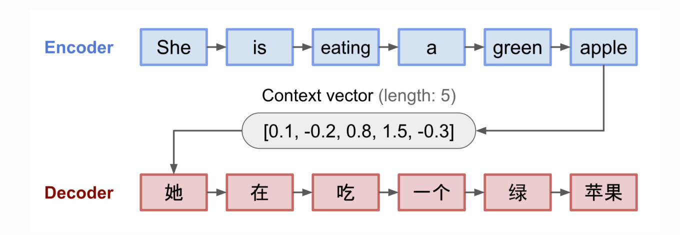

Figure: The encoder-decoder model for translation. (source:https://lilianweng.github.io/lil-log/2018/06/24/attention-attention.html)

- Different parts of input related to different parts of output.
- Encoding complete content difficult (even for LSTMs).
- **Issue:** context vector h_T provides no access to earlier inputs!

WAHT IS ATTENTION

- Key idea: Allow the decoder to access all the hidden states of the encoder (instead of just the final one) so that it can dynamically decide which ones are relevant at each time-step in the decoding.
- This means the decoder can choose to "focus" on different hidden states (of the encoder) at different time-steps of the decoding process similar to how the human eye can focus on different regions of the visual field.
- This is known as an **attention mechanism**.

WAHT IS ATTENTION

- The attention mechanism is implemented by an additional component in the decoder.
- For example, this can be a simple single-hidden layer feed-forward neural network which is trained along with the RNN.
- At any given time-step i of the decoding process, the network computes the relevance of encoder state $\mathbf{z}^{[j]}$ as:

$$rel(\mathbf{z}^{[j]})^{[i]} = \mathbf{v}_a^\top \tanh(\mathbf{W}_a[\mathbf{g}^{[i-1]}; \mathbf{z}^{[j]}])$$

  where $\mathbf{v}_a$ and $\mathbf{W}_a$ are the parameters of the feed-forward network, $\mathbf{g}^{[i-1]}$ is the decoder state from the previous time-step and ';' indicates concatenation.
- The relevance scores (for all the encoder hidden states) are then normalized which gives the *attention weights* $(\alpha^{[j]})^{[i]}$:

$$(\alpha^{[j]})^{[i]} = \frac{\exp(rel(\mathbf{z}^{[j]})^{[i]})}{\sum_{j'} \exp(rel(\mathbf{z}^{[j']})^{[i]})}$$

WAHT IS ATTENTION

- The attention mechanism allows the decoder network to focus on different parts of the input sequence by adding connections from all hidden states of the encoder to each hidden state of the decoder.

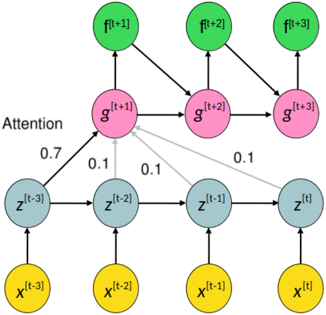

Figure: Attention at $i = t + 1$

WAHT IS ATTENTION

- At each time step *i*, a set of weights $(\alpha^{[j]})^{[i]}$ is computed which determine how to combine the hidden states of the encoder into a context vector $\mathbf{g}^{[i]} = \sum_{j=1}^{n_x} (\alpha^{[j]})^{[i]} \mathbf{z}^{[j]}$, which holds the necessary information to predict the correct output.
- Each hidden state contains mostly information from recent inputs. In the case of a bidirectional RNN to encode the input sequence, a hidden state contains information from recent preceding and following inputs.

WAHT IS ATTENTION

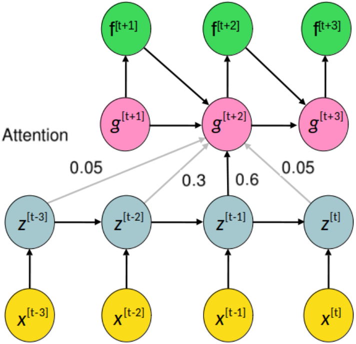

Figure: Attention at $i = t + 2$

WAHT IS ATTENTION

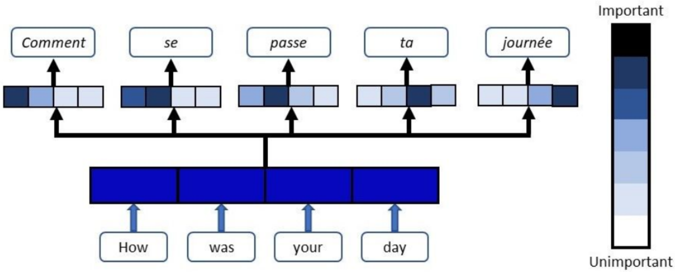

Credit: Gabriel Loye

Figure: An illustration of a machine translation task using an encoder-decoder model with an attention mechanism. The attention weights at each time-step of the decoding/translation process indicate which parts of the input sequence are most relevant. (There are 4 attention weights because there are 4 encoder states.)

ATTENTION

Figure: Attention for image captioning: the attention mechanism tells the network roughly which pixels to pay attention to when writing the text (Kelvin Xu al. 2015)

Transformers

TRANSFORMERS

- Advanced RNNs have similar limitations as vanilla RNN networks:
  - RNNs process the input data sequentially.
  - Difficulties in learning long term dependency (although GRU or LSTM perform better than vanilla RNNs, they sometimes struggle to remember the context introduced earlier in long sequences).
- These challenges are tackled by transformer networks.

TRANSFORMERS

- Transformers are solely based on attention (no RNN or CNN).
- In fact, the paper which coined the term *transformer* is called *Attention is all you need*.
- They are the state-of-the-art networks in natural language processing (NLP) tasks since 2017.
- Transformer architectures like BERT (Bidirectional Encoder Representations from Transformers, 2018) and GPT-3 (Generative Pre-trained Transformer-3, 2020) are pre-trained on a large corpus and can be fine-tuned to specific language tasks.

TRANSFORMERS

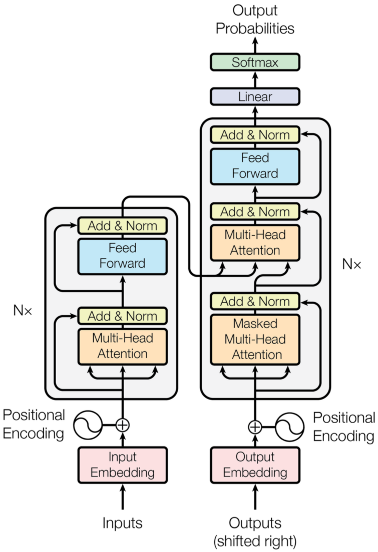

Figure 1: The Transformer - model architecture.

CNNs or RNNs?

CNNS OR RNNS?

- Historically, RNNs were the default for sequence processing tasks.
- However, some families of CNNs (especially those based on Fully Convolutional Networks (FCNs)) *can* be used to process variable-length sequences such as text or time-series data.
- If a CNN doesn't contain any fully-connected layers, the total number of weights in the network is independent of the spatial dimensions of the input because of weight-sharing in the convolutional layers.
- Recent research [Bai et al. , 2018] indicates that such convolutional architectures, so-called Temporal Convolutional Networks (TCNs), can outperform RNNs on a wide range of tasks.
- A major advantage of TCNs is that the entire input sequence can be fed to the network at once (as opposed to sequentially).

CNNS OR RNNS?

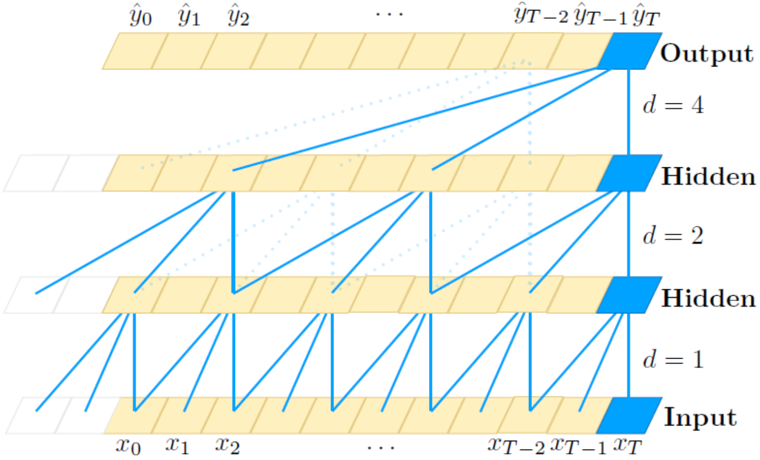

Figure: A TCN (we have already seen this in the CNN lecture!) is simply a variant of the one-dimensional FCN which uses a special type of dilated convolutions called **causal dilated** convolutions.

SUMMARY

- RNNs are specifically designed to process sequences of varying lengths.
- For that recurrent connections are introduced into the network structure.
- The gradient is calculated by backpropagation through time.
- An LSTM replaces the simple hidden neuron by a complex system consisting of cell state, and forget, input, and output gates.
- An RNN can be used as a language model, which can be improved by word-embeddings.
- Different advanced types of RNNs exist, like Encoder-Decoder architectures and bidirectional RNNs.[1]

1. A bidirectional RNN processes the input sequence in both directions (front-to-back and back-to-front).

REFERENCES

Ian Goodfellow, Yoshua Bengio and Aaron Courville (2016)
Deep Learning
http://www.deeplearningbook.org/

Oriol Vinyals, Alexander Toshev, Samy Bengio and Dumitru Erhan (2014)
Show and Tell: A Neural Image Caption Generator
https://arxiv.org/abs/1411.4555

Alex Graves (2013)
Generating Sequences With Recurrent Neural Networks
https://arxiv.org/abs/1308.0850

Namrata Anand and Prateek Verma (2016)
Convolutional and recurrent nets for detecting emotion from audio data
http://cs231n.stanford.edu/reports/2015/pdfs/Cs_231n_paper.pdf

Gabriel Loye (2019)
Attention Mechanism
https://blog.floydhub.com/attention-mechanism/

REFERENCES

Andrew Owens, Phillip Isola, Josh H. McDermott, Antonio Torralba, Edward H. Adelson and William T. Freeman (2015)
Visually Indicated Sounds
https://arxiv.org/abs/1512.08512

Andrej Karpathy (2015)
The Unreasonable Effectiveness of Recurrent Neural Networks
http://karpathy.github.io/2015/05/21/rnn-effectiveness/

Kelvin Xu, Jimmy Ba, Ryan Kiros, Kyunghyun Cho, Aaron C. Courville, Ruslan Salakhutdinov, Richard S. Zemel and Yoshua Bengio (2015)
Show, Attend and Tell: Neural Image Caption Generation with Visual Attention
https://arxiv.org/abs/1502.03044

Shaojie Bai, J. Zico Kolter, Vladlen Koltun (2018)
An Empirical Evaluation of Generic Convolutional and Recurrent Networks for Sequence Modeling
https://arxiv.org/abs/1803.01271

REFERENCES

Lilian Weng (2018)
Attention? Attention!
https://lilianweng.github.io/lil-log/2018/06/24/attention-attention.html