


# Deep Learning

## Autoencoders - Basic Principle



### Learning goals

- Task and structure of an AE
- Undercomplete AEs
- Relation of AEs and PCA

# AUTOENCODER-TASK AND STRUCTURE


- Autoencoders (AEs) are NNs for unsupervised learning of a lower dimensional feature representation from unlabeled training data.
- Task: Learn a compression of the data.
- Autoencoders consist of two parts:
  - **encoder** learns mapping from the data  $\mathbf{x}$  to a low-dimensional latent variable  $\mathbf{z} = \text{enc}(\mathbf{x})$ .
  - **decoder** learns mapping back from latent  $\mathbf{z}$  to a reconstruction  $\hat{\mathbf{x}} = \text{dec}(\mathbf{z})$  of  $\mathbf{x}$ .
- Loss function does not use any labels and measures the quality of the reconstruction compared to the input:

$$L(\mathbf{x}, \text{dec}(\text{enc}(\mathbf{x})))$$

- Goal: Learn good **representation  $\mathbf{z}$**  (also called **code**).

# AUTOENCODER (AE)- COMPUTATIONAL GRAPH

The general structure of an AE as a computational graph:




- An AE has two computational steps:
  - the encoder  $enc$ , mapping  $\mathbf{x}$  to  $\mathbf{z}$ .
  - the decoder  $dec$ , mapping  $\mathbf{z}$  to  $\hat{\mathbf{x}}$ .

# Undercomplete Autoencoders


# UNDERCOMPLETE AUTOENCODERS

- A naive implementation of an autoencoder would simply learn the identity  $dec(enc(\mathbf{x})) = \mathbf{x}$ .
- This would not be useful.



# UNDERCOMPLETE AUTOENCODERS

- A naive implementation of an autoencoder would simply learn the identity  $dec(enc(\mathbf{x})) = \mathbf{x}$ .
- This would not be useful.




# UNDERCOMPLETE AUTOENCODERS

- Therefore we have a “bottleneck” layer: We restrict the architecture, such that


$$\dim(\mathbf{z}) < \dim(\mathbf{x})$$

- Such an AE is called **undercomplete**.



# UNDERCOMPLETE AUTOENCODERS

- In an undercomplete AE, the hidden layer has fewer neurons than the input layer.
- That will force the AE to
- capture only the most salient features of the training data!
  - learn a “compressed” representation of the input.





# UNDERCOMPLETE AUTOENCODERS

- Training an AE is done by minimizing the risk with a loss function penalizing the reconstruction  $dec(enc(\mathbf{x}))$  for differing from  $\mathbf{x}$ .
- The L2-loss


$$\|\mathbf{x} - dec(enc(\mathbf{x}))\|_2^2$$

is a typical choice, but other loss functions are possible.

- For optimization, the same optimization techniques as for standard feed-forward nets are applied (SGD, RMSProp, ADAM,...).


# EXPERIMENT: LEARN TO ENCODE MNIST

- Let us try to compress the MNIST data as good as possible.
- We train undercomplete AEs with different dimensions of the internal representation  $\mathbf{z}$  (i.e. different “bottleneck” sizes).




**Figure:** Flow chart of our autoencoder: reconstruct the input with fixed dimensions  $\dim(\mathbf{z}) \leq \dim(\mathbf{x})$ .

# EXPERIMENT: LEARN TO ENCODE MNIST



**Figure:** Architecture of the autoencoder.


# EXPERIMENT: LEARN TO ENCODE MNIST



**Figure:** The top row shows the original digits, the bottom row the reconstructed ones.

- $\dim(\mathbf{z}) = 784 = \dim(\mathbf{x})$ .


# EXPERIMENT: LEARN TO ENCODE MNIST



**Figure:** The top row shows the original digits, the bottom row the reconstructed ones.

- $\dim(\mathbf{z}) = 256$ .

# EXPERIMENT: LEARN TO ENCODE MNIST



**Figure:** The top row shows the original digits, the bottom row the reconstructed ones.

- $\dim(\mathbf{z}) = 64$ .


# EXPERIMENT: LEARN TO ENCODE MNIST



**Figure:** The top row shows the original digits, the bottom row the reconstructed ones.

- $\dim(\mathbf{z}) = 32$ .

# EXPERIMENT: LEARN TO ENCODE MNIST




**Figure:** The top row shows the original digits, the bottom row the reconstructed ones.

- $\dim(\mathbf{z}) = 16$ .




# EXPERIMENT: LEARN TO ENCODE MNIST



**Figure:** The top row shows the original digits, the bottom row the reconstructed ones.

- $\dim(\mathbf{z}) = 8$ .


# EXPERIMENT: LEARN TO ENCODE MNIST



**Figure:** The top row shows the original digits, the bottom row the reconstructed ones.

- $\dim(\mathbf{z}) = 4$ .


# EXPERIMENT: LEARN TO ENCODE MNIST



**Figure:** The top row shows the original digits, the bottom row the reconstructed ones.

- $\dim(\mathbf{z}) = 2$ .

# EXPERIMENT: LEARN TO ENCODE MNIST




**Figure:** The top row shows the original digits, the bottom row the reconstructed ones.

- $\dim(\mathbf{z}) = 1$ .

# INCREASING THE CAPACITY OF AES


Increasing the number of layers adds capacity to autoencoders:



# **Autoencoders as Principal Component Analysis**


# AES AS PRINCIPAL COMPONENT ANALYSIS

- Consider a undercomplete autoencoder with
  - **linear** encoder function  $enc(\mathbf{x})$ , and
  - **linear** decoder function  $dec(\mathbf{z})$ .
- The L2-loss  $\|\mathbf{x} - dec(enc(\mathbf{x}))\|_2^2$  is employed and inputs are normalized to zero mean.
- We want to find the **linear projection** of the data with the minimal L2-reconstruction error.



# AES AS PRINCIPAL COMPONENT ANALYSIS


- It can be shown that the optimal solution is an **orthogonal** linear transformation (i.e. a rotation of the coordinate system) given by the  $\dim(\mathbf{z}) = k$  singular vectors with largest singular values.





# AES AS PRINCIPAL COMPONENT ANALYSIS


- This is an equivalent formulation to **Principal Component Analysis (PCA)**, which uses an orthogonal transformation to convert a set of observations of possibly correlated variables into a set of values of linearly uncorrelated variables called **principal components**.
- The transformation is defined in such a way that the first principal component has the largest possible variance (i.e., accounts for as much of the variability in the data as possible).



# AES AS PRINCIPAL COMPONENT ANALYSIS

- The formulations are equivalent: “Find a linear projection into a  $k$ -dimensional space that ...”
  - “... minimizes the L2-reconstruction error” (AE-based formulation).
  - “... maximizes the variance of the projected datapoints” (statistical formulation).
- An AE with a non-linear decoder/encoder can be seen as a non-linear generalization of PCA.

Linear vs nonlinear dimensionality reduction



**Figure:** Credits: Jeremy Jordan “Introduction to autoencoders”

# REFERENCES



Ian Goodfellow, Yoshua Bengio and Aaron Courville (2016)

Deep Learning

*<http://www.deeplearningbook.org/>*