

# Deep Learning

## Convolutional Operation

a	b	c
d	e	f
g	h	i

Input: 3x3x1

$w_{11}$	$w_{12}$
$w_{21}$	$w_{22}$

Filter: 2x2x1

$s_{11}$	$s_{12}$
$s_{21}$	$s_{22}$

Output: 2x2x1

$$s_{11} = a \cdot w_{11} + b \cdot w_{12} + d \cdot w_{21} + e \cdot w_{22}$$

$$s_{12} = b \cdot w_{11} + c \cdot w_{12} + e \cdot w_{21} + f \cdot w_{22}$$

$$s_{21} = d \cdot w_{11} + e \cdot w_{12} + g \cdot w_{21} + h \cdot w_{22}$$

$$s_{22} = e \cdot w_{11} + f \cdot w_{12} + h \cdot w_{21} + i \cdot w_{22}$$

### Learning goals

- What are filters?
- Convolutional Operation
- 2D Convolution

# FILTERS TO EXTRACT FEATURES

- Filters are widely applied in Computer Vision (CV) since the 70's.
- One prominent example: **Sobel-Filter**.
- It detects edges in images.



**Figure:** Sobel-filtered image.

# FILTERS TO EXTRACT FEATURES

- Edges occur where the intensity over neighboring pixels changes fast.
- Thus, approximate the gradient of the intensity of each pixel.
- Sobel showed that the gradient image  $\mathbf{G}_x$  of original image  $\mathbf{A}$  in x-dimension can be approximated by:

$$\mathbf{G}_x = \begin{bmatrix} -1 & 0 & +1 \\ -2 & 0 & +2 \\ -1 & 0 & +1 \end{bmatrix} * \mathbf{A} = \mathbf{S}_x * \mathbf{A}$$

where  $*$  indicates a mathematical operation known as a **convolution**, not a traditional matrix multiplication.

- The filter matrix  $\mathbf{S}_x$  consists of the product of an **averaging** and a **differentiation** kernel:

$$\underbrace{\begin{bmatrix} 1 & 2 & 1 \end{bmatrix}}_{\text{averaging}}^T \underbrace{\begin{bmatrix} -1 & 0 & +1 \end{bmatrix}}_{\text{differentiation}}$$

# FILTERS TO EXTRACT FEATURES

- Similarly, the gradient image  $\mathbf{G}_y$  in y-dimension can be approximated by:

$$\mathbf{G}_y = \begin{bmatrix} -1 & -2 & -1 \\ 0 & 0 & 0 \\ +1 & +2 & +1 \end{bmatrix} * \mathbf{A} = \mathbf{S}_y * \mathbf{A}$$

- The combination of both gradient images yields a dimension-independent gradient information  $\mathbf{G}$ :

$$\mathbf{G} = \sqrt{\mathbf{G}_x^2 + \mathbf{G}_y^2}$$

- These matrix operations were used to create the filtered picture of Albert Einstein.

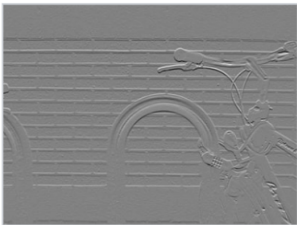
# HORIZONTAL VS VERTICAL EDGES



Input



Vertical edges detected by  $S_x$



Horizontal edges detected by  $S_y$

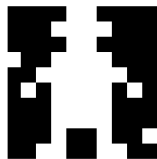


Combined

Source: Wikipedia

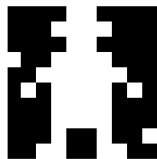
**Figure:** Sobel filtered images. Outputs are normalized in each case.

# FILTERS TO EXTRACT FEATURES



- Let's do this on a dummy image.
- How to represent a digital image?

# FILTERS TO EXTRACT FEATURES



0	0	0	0	255	255	0	0	0	0
0	0	0	255	255	255	255	0	0	0
0	0	0	0	255	255	0	0	0	0
255	0	0	255	255	255	255	0	0	0
0	0	255	255	255	255	255	255	0	0
0	255	0	255	255	255	255	0	255	0
0	0	0	255	255	255	255	0	0	0
0	0	0	255	255	255	255	0	0	0
0	0	0	255	255	0	0	255	0	255
0	0	255	255	0	0	255	255	0	0

- Basically as an array of integers.

# FILTERS TO EXTRACT FEATURES

Sobel-Operator

$$S_x = \begin{bmatrix} -1 & 0 & +1 \\ -2 & \boxed{0} & +2 \\ -1 & 0 & +1 \end{bmatrix}$$

0	0	0	0	255	255	0	0	0	0
0	0	0	255	255	255	255	0	0	0
0	0	0	0	255	255	0	0	0	0
255	0	0	255	255	255	255	0	0	0
0	0	255	255	255	255	255	255	0	0
0	255	0	255	255	255	255	0	255	0
0	0	0	255	255	255	255	0	0	0
0	0	0	255	255	255	255	0	0	0
0	0	0	255	0	0	255	0	0	255
0	0	255	255	0	0	255	255	0	0

- $S_x$  enables us to detect vertical edges!



# FILTERS TO EXTRACT FEATURES

Sobel-Operator

$$S_x = \begin{bmatrix} -1 & 0 & +1 \\ -2 & 0 & +2 \\ -1 & 0 & +1 \end{bmatrix}$$

0	0	0	0	255	255	0	0	0	0
0	0	0	255	255	255	255	0	0	0
0	0	0	0	0	255	255	0	0	0
255	0	0	255	255	255	255	255	0	0
0	0	255	255	255	255	255	255	0	0
0	255	0	255	255	255	255	0	255	0
0	0	0	255	255	255	255	0	0	0
0	0	0	255	255	255	255	0	0	0
0	0	0	255	0	0	255	0	0	255
0	0	255	255	0	0	255	255	0	0

# FILTERS TO EXTRACT FEATURES

Sobel-Operator

$$S_x = \begin{bmatrix} -1 & 0 & +1 \\ -2 & \boxed{0} & +2 \\ -1 & 0 & +1 \end{bmatrix}$$

0	0	0	0	255	255	0	0	0	0
0	0	0	255	255	255	255	0	0	0
0	0	0	<span style="border: 1px solid red;">0</span>	255	255	255	0	0	0
255	0	0	255	255	255	255	255	0	0
0	0	255	255	255	255	255	255	0	0
0	255	0	255	255	255	255	0	255	0
0	0	0	255	255	255	255	0	0	0
0	0	0	255	255	255	255	0	0	0
0	0	0	255	0	0	255	0	0	255
0	0	255	255	0	0	255	255	0	0

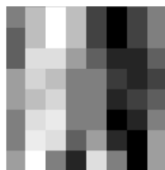
$$\begin{aligned} (G_x)_{(i,j)} &= (I \star S_x)_{(i,j)} = -1 \cdot 0 + 0 \cdot 255 + \mathbf{1 \cdot 255} \\ &\quad - 2 \cdot 0 + 0 \cdot 0 + \mathbf{2 \cdot 255} \\ &\quad - 1 \cdot 0 + 0 \cdot 255 + \mathbf{1 \cdot 255} \end{aligned}$$

# FILTERS TO EXTRACT FEATURES

0	510	1020	510	-510	-1020	-510	0
-255	510	1020	510	-510	-1020	-510	0
-255	765	765	255	-255	-765	-765	-255
255	765	510	0	0	-510	-765	-510
255	510	765	0	0	-765	-510	-255
0	765	1020	0	0	-1020	-765	0
0	1020	765	-255	255	-765	-1020	255
255	1020	0	-765	765	0	-1020	255

- Applying the Sobel-Operator to every location in the input yields the **feature map**.

# FILTERS TO EXTRACT FEATURES



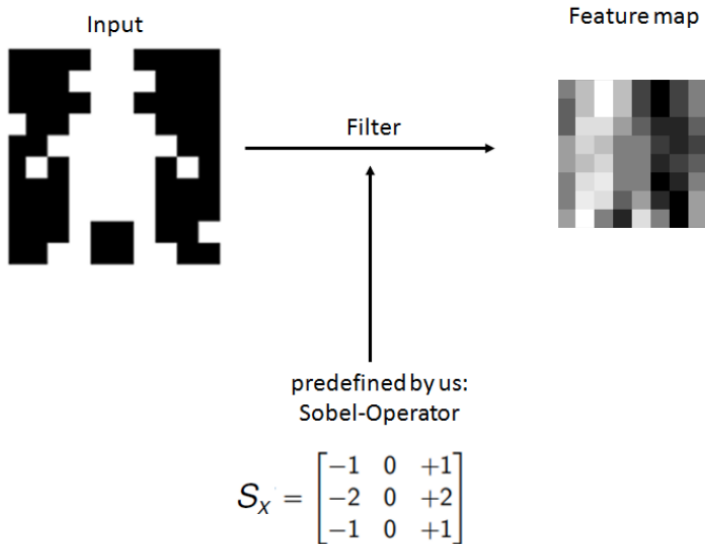
128	191	255	191	64	0	64	128
96	191	255	191	64	0	64	128
96	223	223	159	96	32	32	96
159	223	191	128	128	64	32	64
159	191	223	128	128	32	64	96
128	223	255	128	128	0	32	128
128	255	223	96	159	32	0	159
159	255	128	32	223	128	0	159

- Normalized feature map reveals vertical edges.
- Note the dimensional reduction compared to the dummy image.

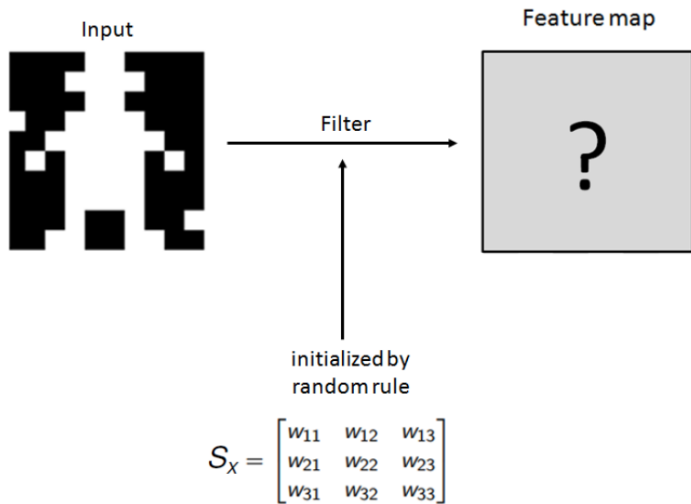
# WHY DO WE NEED TO KNOW ALL OF THAT?

- What we just did was extracting **pre-defined** features from our input (i.e. edges).
- A convolutional neural network does almost exactly the same: “extracting features from the input”.  
⇒ The main difference is that we usually do not tell the CNN what to look for (pre-define them), **the CNN decides itself**.
- In a nutshell:
  - We initialize a lot of random filters (like the Sobel but just random entries) and apply them to our input.
  - Then, a classifier which (e.g. a feed forward neural net) uses them as input data.
  - Filter entries will be adjusted by common gradient descent methods.

# WHY DO WE NEED TO KNOW ALL OF THAT?



# WHY DO WE NEED TO KNOW ALL OF THAT?



# WORKING WITH IMAGES

- In order to understand the functionality of CNNs, we have to familiarize ourselves with some properties of images.
- Grey scale images:
  - Matrix with dimensions **height**  $\times$  **width**  $\times$  1.
  - Pixel entries differ from 0 (black) to 255 (white).
- Color images:
  - Tensor with dimensions **height**  $\times$  **width**  $\times$  3.
  - The depth 3 denotes the RGB values (red - green - blue).
- Filters:
  - A filter's depth is **always** equal to the input's depth!
  - In practice, filters are usually square.
  - Thus we only need one integer to define its size.
  - For example, a filter of size 2 applied on a color image actually has the dimensions  $2 \times 2 \times 3$ .



# THE 2D CONVOLUTION

- Suppose we have an input with entries  $a, b, \dots, i$  (think of pixel values).
- The filter we would like to apply has weights  $w_{11}, w_{12}, w_{21}$  and  $w_{22}$ .

a	b	c
d	e	f
g	h	i

Input: 3x3x1

$w_{11}$	$w_{12}$
$w_{21}$	$w_{22}$

Filter: 2x2x1

# THE 2D CONVOLUTION

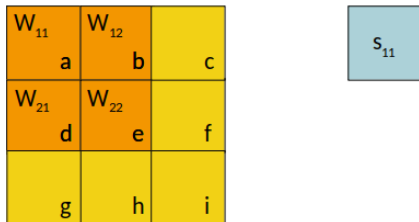
- Suppose we have an input with entries  $a, b, \dots, i$  (think of pixel values).
- The filter we would like to apply has weights  $w_{11}, w_{12}, w_{21}$  and  $w_{22}$ .

$w_{11}$ a	$w_{12}$ b	c
$w_{21}$ d	$w_{22}$ e	f
g	h	i



# THE 2D CONVOLUTION

- Suppose we have an input with entries  $a, b, \dots, i$  (think of pixel values).
- The filter we would like to apply has weights  $w_{11}, w_{12}, w_{21}$  and  $w_{22}$ .



To obtain  $s_{11}$  we simply compute the dot product:

$$s_{11} = a \cdot w_{11} + b \cdot w_{12} + d \cdot w_{21} + e \cdot w_{22}$$

# THE 2D CONVOLUTION

- Suppose we have an input with entries  $a, b, \dots, i$  (think of pixel values).
- The filter we would like to apply has weights  $w_{11}, w_{12}, w_{21}$  and  $w_{22}$ .

	$w_{11}$	$w_{12}$
$a$	$b$	$c$
$d$	$e$	$f$
$g$	$h$	$i$

$s_{11}$	$s_{12}$
----------	----------

Same for  $s_{12}$ :

$$s_{12} = b \cdot w_{11} + c \cdot w_{12} + e \cdot w_{21} + f \cdot w_{22}$$

# THE 2D CONVOLUTION

- Suppose we have an input with entries  $a, b, \dots, i$  (think of pixel values).
- The filter we would like to apply has weights  $w_{11}, w_{12}, w_{21}$  and  $w_{22}$ .

	a	b	c
$w_{11}$	d	$w_{12}$	e
$w_{21}$	g	$w_{22}$	h

$s_{11}$	$s_{12}$
$s_{21}$	

As well as for  $s_{21}$ :

$$s_{21} = d \cdot w_{11} + e \cdot w_{12} + g \cdot w_{21} + h \cdot w_{22}$$

# THE 2D CONVOLUTION

- Suppose we have an input with entries  $a, b, \dots, i$  (think of pixel values).
- The filter we would like to apply has weights  $w_{11}, w_{12}, w_{21}$  and  $w_{22}$ .

a	b	c
d	$w_{11}$ e	$w_{12}$ f
g	$w_{21}$ h	$w_{22}$ i

$s_{11}$	$s_{12}$
$s_{21}$	$s_{22}$

And finally for  $s_{22}$ :

$$s_{22} = e \cdot w_{11} + f \cdot w_{12} + h \cdot w_{21} + i \cdot w_{22}$$

# THE 2D CONVOLUTION

- Suppose we have an input with entries  $a, b, \dots, i$  (think of pixel values).
- The filter we would like to apply has weights  $w_{11}, w_{12}, w_{21}$  and  $w_{22}$ .

a	b	c
d	e	f
g	h	i

Input: 3x3x1

$w_{11}$	$w_{12}$
$w_{21}$	$w_{22}$

Filter: 2x2x1

$s_{11}$	$s_{12}$
$s_{21}$	$s_{22}$

Output: 2x2x1

$$s_{11} = a \cdot w_{11} + b \cdot w_{12} + d \cdot w_{21} + e \cdot w_{22}$$

$$s_{12} = b \cdot w_{11} + c \cdot w_{12} + e \cdot w_{21} + f \cdot w_{22}$$

$$s_{21} = d \cdot w_{11} + e \cdot w_{12} + g \cdot w_{21} + h \cdot w_{22}$$

$$s_{22} = e \cdot w_{11} + f \cdot w_{12} + h \cdot w_{21} + i \cdot w_{22}$$

# THE 2D CONVOLUTION

- Suppose we have an input with entries  $a, b, \dots, i$  (think of pixel values).
- The filter we would like to apply has weights  $w_{11}, w_{12}, w_{21}$  and  $w_{22}$ .

a	b	c
d	e	f
g	h	i

Input: 3x3x1

$w_{11}$	$w_{12}$
$w_{21}$	$w_{22}$

Filter: 2x2x1

$s_{11}$	$s_{12}$
$s_{21}$	$s_{22}$

Output: 2x2x1

More generally, let  $I$  be the matrix representing the input and  $W$  be the filter/kernel. Then the entries of the output matrix are defined by  $s_{ij} = \sum_{m,n} I_{i+m-1,j+n-1} w_{mn}$  where  $m, n$  denote the image size and kernel size respectively.